

k-Nearest Neighbors (kNN)

Chris Reger

Original Slides: Ryan Henning, Brandon Martin-Anderson, Lori Bordzuk



Today's Objectives

- Implement KNN algorithm
- Explain the difference between KNN for regression vs. classification
- Understand KNN hyperparameters
 - How does changing them affect the model?
- Describe the curse of dimensionality

Supervised vs. Unsupervised Learning

Supervised learning is the machine learning task of learning a function that maps an input to an output based on example input-output pairs.

X, y -> predicting y based on the values in X

X, y a.k.a

features, target

independent, dependent

exogenous, endogenous

predictors, response

Example capstone: [Avalanche Prediction](#)

Unsupervised learning is a type of self-organized ... learning that helps find previously unknown patterns in data set without pre-existing labels.

X -> understanding structure in X

Example capstone: [Spice Blends](#)

Parametric vs Non-parametric models

A machine learning algorithm can be supervised or unsupervised, and parametric or non-parametric.


A **parametric** algorithm

- has a fixed number of parameters
- makes assumptions about the structure of the data
- will work well if the assumptions are correct!
- common examples: linear regression, neural networks, statistical distributions defined by a finite set of parameters

A **non-parametric** algorithm

- uses a flexible number of parameters, and the number of parameters often grows as it learns from more data.
- makes fewer assumptions about the data
- common examples: K-Nearest Neighbors, decision trees

Big dog or small horse?



New datapoint 1:
Is it a dog or a horse?

Horse

New datapoint 3:
Is it a dog or a horse?

Great Dane?

Pony?

New datapoint 2:
Is it a dog or a horse?

Dog


The k-Nearest Neighbors: Classification

For a new input x , predict the most common label amongst its k closest neighbors

Image on right:

$k = 3$

Predict **BLUE**




The k-Nearest Neighbors: Regression

For a new input x , predict the **average label** amongst its k closest neighbors

Image on right:

$k = 3$

Predict **13.3**



The k-Nearest Neighbors Algorithm

Training algorithm:

1. Store all the data.

Prediction algorithm (predict the class of a new point x'):


1. Calculate the distance from x' to all points in your dataset.
2. Sort the points in your dataset by increasing distance from x' .
3. Predict the majority label of the k closest points.

kNN Hyperparameter: Distance Metrics

Euclidean Distance (L2):
$$\sum_i (a_i - b_i)^2$$

Manhattan Distance (L1):
$$\sum_i |a_i - b_i|$$

Cosine Distance = 1 - Cosine Similarity:
$$1 - \frac{a \cdot b}{||a|| ||b||}$$



What is the prediction
when $k=3$?


Class B

What is the prediction
when $k=6$?


Class A

Hyperparameter k : the number of nearest neighbors to consider

$k=1$




$k=5$




Hyperparameter k : the number of nearest neighbors to consider

$k=10$




$k=50$



Which model is overfit?


k=1



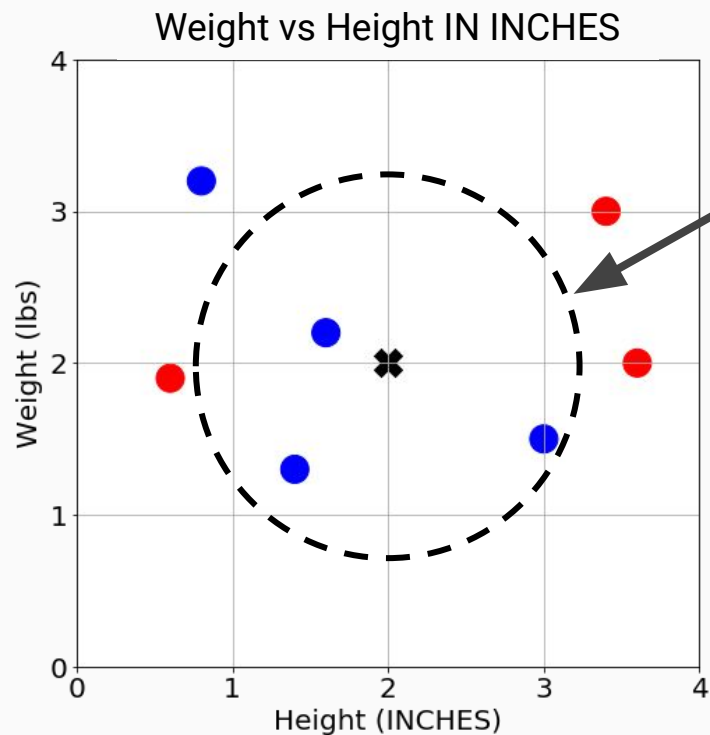
As a
general
rule, start
with:

$$k = \sqrt{n}$$

k=50



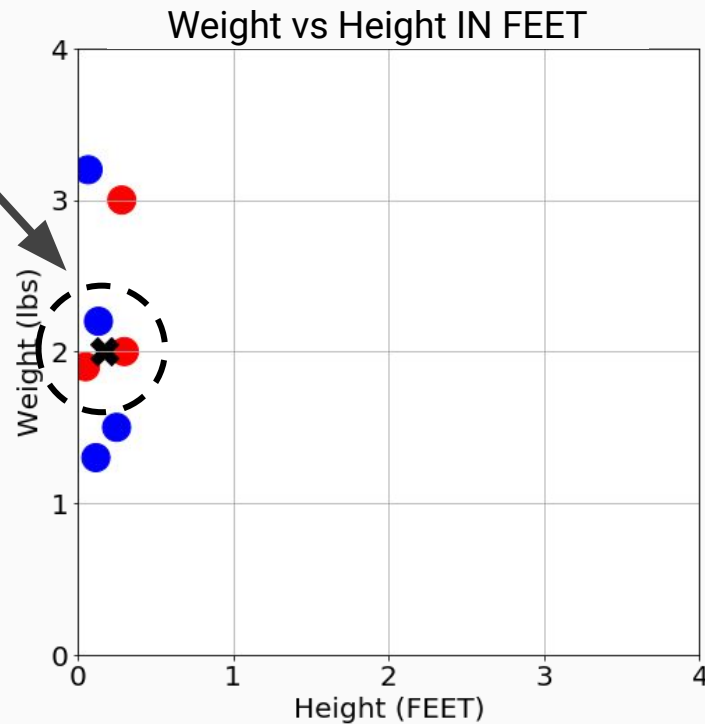
Be careful with the scale of your features!




$k = 3$

The three
“closest
neighbors”
differ
depending on
the scale of the
feature....

**Don't forget
to scale your
data!!!!**



$k = 3$



Let the k nearest points have distances:

$$d_1, d_2, \dots, d_k$$

The i^{th} point votes with a weight of:



$$\frac{1}{d_i}$$

small distances are weighted more!

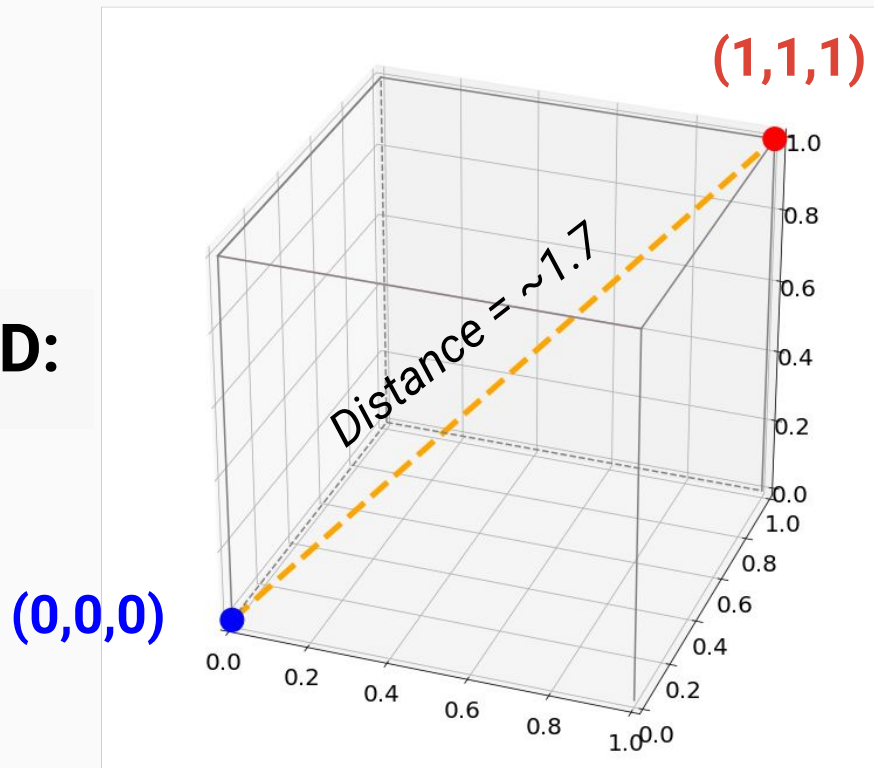
kNN in high dimensions

kNN works pretty well (in *general*) for dimensions < 5
but is problematic when used with high dimensional spaces

In high dimensions, the nearest neighbors can be very “far away”




3D:



The Curse of Dimensionality (another perspective)


Given 100 (random) sample points....



1D

All 100 pts within 1 unit
of center.


Density = 100



2D

77 pts within 1 unit of
center.

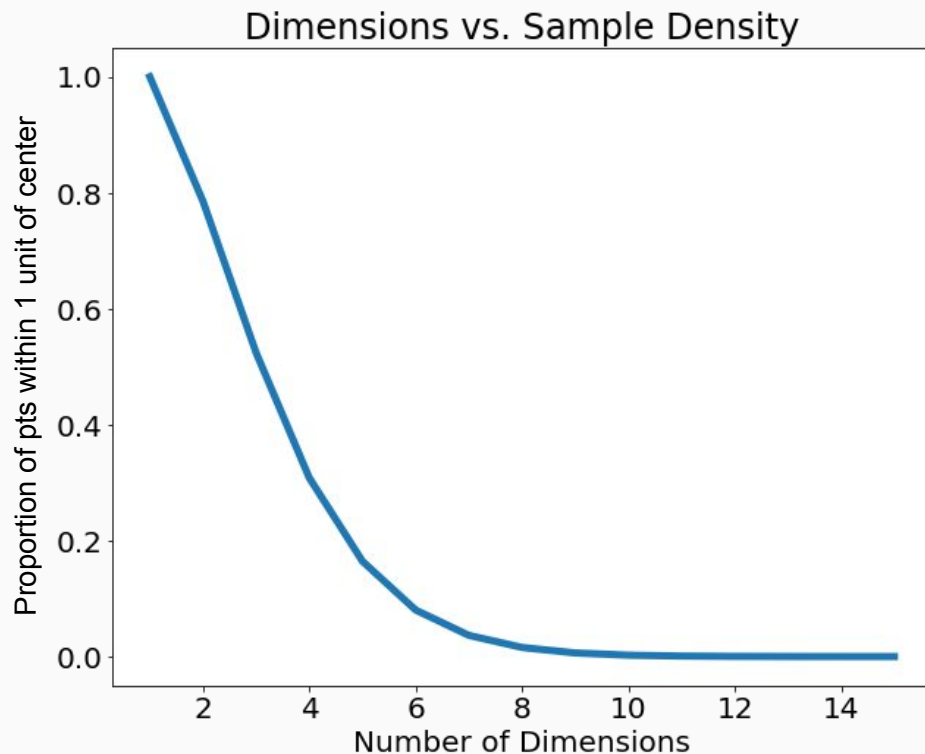
Density = 77



3D

53 pts within 1 unit of
center.

Density = 53



The **more dimensions** you have, the **more data points** you need to maintain density.

General guideline:

- Given n data points in d_{orig} dimensions...
- If you want to *increase* the total number of dimensions to d_{new} , you now need:
- $n \frac{d_{new}}{d_{orig}}$ data points to maintain density

The Curse of Dimensionality takeaways

- kNN (or any method that relies on distance metrics) will suffer in high dimensions.
 - Nearest neighbors are “far” away in high dimensions (even for $d=10$).
- High dimensional data tends to be sparse; it's easy to overfit sparse data.
 - It takes A LOT OF DATA to make up for increased dimensionality.

Summary: kNN

Pros:

- Super simple
- Training is trivial (store the data)
- Works with any number of classes
- Easy to add more data
- Few hyperparameters:
 - *distance metric*
 - *k*

Cons:

- High prediction cost (especially for large datasets)
- Bad with high dimensions
 - you'll learn dimensionality reduction methods later on!
- Categorical features don't work well

Review: Today's Objectives

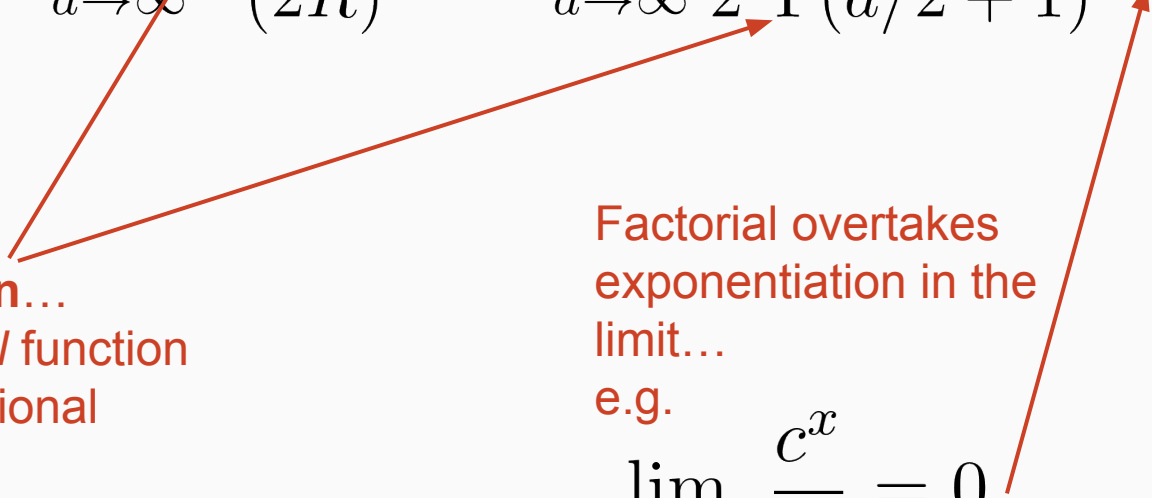
- Implement KNN algorithm
- Explain the difference between KNN for regression vs. classification
- Understand KNN hyperparameters
 - How does changing them affect the model?
- Describe the curse of dimensionality

Morning Exercise: Implement a sklearn-style KNN algorithm

Appendix

 galvanize

Don't freak out...

$$\lim_{d \rightarrow \infty} \frac{V_{\text{sphere}}(R, d)}{V_{\text{cube}}(R, d)} = \lim_{d \rightarrow \infty} \frac{\frac{\pi^{d/2} R^d}{\Gamma(d/2 + 1)}}{(2R)^d} = \lim_{d \rightarrow \infty} \frac{\pi^{d/2}}{2^d \Gamma(d/2 + 1)} = 0$$


Euler's gamma function...

basically, it's the *factorial* function that can operate on fractional numbers

What does this mean?

Factorial overtakes exponentiation in the limit...

e.g.

$$\lim_{x \rightarrow \infty} \frac{c^x}{x!} = 0$$

Parametric vs Non-parametric Models

Parametric models have a fixed number of learned parameters.

- Logistic regression is parametric.
- kNN is non-parametric.

Parametric models are more structured. The added structure often combats the curse of dimensionality... as long as the structure is derived from reasonable assumptions.

Alternate perspective: Parametric models are not distance based, so the curse doesn't apply!